

# Movie Rating Classification

## A Data Science Workflow

**IMDb Most Popular Movies (2006-2016)**

### *Project Write-Up*

#### **Project Overview:**

This project applies a comprehensive data science workflow to analyze and classify movies based on their ratings using Python. The goal is to understand what factors influence movie ratings and build a classification model to predict whether a movie will be highly rated or not.

**Dataset:** IMDb Most Popular Movies (2006-2016)

**Source:** <https://www.kaggle.com/datasets/PromptCloudHQ/imdb-data>

## Table of Contents

<b>Step 1:</b>	Load & Explore Dataset	3
<b>Step 2:</b>	Data Cleaning	4
<b>Step 3:</b>	Create Target Label	5
<b>Step 4:</b>	Feature Engineering	6
<b>Step 5:</b>	Exploratory Data Analysis	7
<b>Step 6:</b>	Train-Test Split	8
<b>Step 7:</b>	Choose & Train Model	9
<b>Step 8:</b>	Model Evaluation	10
<b>Step 9:</b>	Interpretation & Insights	11

## Step 1: Load & Explore Dataset

The first step in any data science project is to load the dataset and perform initial exploration to understand its structure, size, and content.

### Code Implementation:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# Load the dataset
df = pd.read_csv('IMDB-Movie-Data.csv')

# Display first few rows
print(df.head())

# Check dataset dimensions
print(f'Dataset shape: {df.shape}')
print(f'Number of rows: {df.shape[0]}')
print(f'Number of columns: {df.shape[1]}')

# Examine data types
print(df.dtypes)

# Check for missing values
print(df.isnull().sum())
```

### Key Findings:

<b>Dataset Size:</b>	1000 movies
<b>Features:</b>	12 columns (Rank, Title, Genre, Description, Director, Actors, Year, Runtime, Rating, Votes, Revenue, Metascore)
<b>Data Types:</b>	Mixed (integers, floats, strings)
<b>Missing Values:</b>	Revenue (128), Metascore (64)

## Step 2: Data Cleaning

Data cleaning is crucial for building accurate models. We remove irrelevant features, handle missing values, and ensure data quality.

### Code Implementation:

```
# Drop irrelevant columns
columns_to_drop = ['Rank', 'Title', 'Description', 'Director', 'Actors']
df_clean = df.drop(columns=columns_to_drop)

# Check missing values percentage
missing_pct = (df_clean.isnull().sum() / len(df_clean)) * 100
print(missing_pct[missing_pct > 0])

# Handle missing values
# For Revenue and Metascore, fill with median
df_clean['Revenue (Millions)'].fillna(
    df_clean['Revenue (Millions)'].median(), inplace=True
)
df_clean['Metascore'].fillna(
    df_clean['Metascore'].median(), inplace=True
)

# Verify no missing values remain
print(f'Missing values after cleaning: {df_clean.isnull().sum().sum()}')
```

### Cleaning Steps:

1. **Removed Columns:** Rank, Title, Description, Director, Actors (non-predictive text data)
2. **Missing Values Handling:**
  - Revenue: Filled with median value (12.8% missing)
  - Metascore: Filled with median value (6.4% missing)
3. **Final Dataset:** 1000 rows with 7 clean columns

## Step 3: Create Target Label

To perform classification, we create a binary target variable that categorizes movies as 'High Rated' or 'Low Rated' based on a rating threshold of 7.0.

### Code Implementation:

```
# Create binary label  
df_clean['label'] = (df_clean['Rating'] >= 7).astype(int)  
  
# Check distribution  
print(df_clean['label'].value_counts())  
print(df_clean['label'].value_counts(normalize=True))
```

### Classification Rules:

High Rated (1):	Rating $\geq$ 7.0
Low Rated (0):	Rating $<$ 7.0

### Label Distribution:

High Rated (1):	544 movies (54.4%)
Low Rated (0):	456 movies (45.6%)

The dataset shows a relatively balanced distribution with slightly more high-rated movies (54.4%) compared to low-rated ones (45.6%).

## Step 4: Feature Engineering

Feature engineering transforms raw data into meaningful features that can be used by machine learning models. We encode categorical variables and select relevant numeric features.

### Code Implementation:

```
from sklearn.preprocessing import LabelEncoder

# Handle Genre (multiple genres per movie)
# Extract first genre for simplicity
df_clean['Genre_Primary'] = df_clean['Genre'].apply(
    lambda x: x.split(',') [0].strip()
)

# Encode genre using Label Encoding
le = LabelEncoder()
df_clean['Genre_Encoded'] = le.fit_transform(df_clean['Genre_Primary'])

# Select features for modeling
feature_columns = [
    'Runtime (Minutes)',
    'Votes',
    'Revenue (Millions)',
    'Metascore',
    'Genre_Encoded',
    'Year'
]

x = df_clean[feature_columns]
y = df_clean['label']

print(f'Feature matrix shape: {x.shape}')
print(f'Target variable shape: {y.shape}'')
```

### Selected Features:

Feature	Type	Description
Runtime (Minutes)	Numeric	Movie duration in minutes
Votes	Numeric	Number of user votes on IMDb
Revenue (Millions)	Numeric	Box office revenue in millions USD
Metascore	Numeric	Metacritic score (0-100)
Genre_Encoded	Categorical (Encoded)	Primary genre label-encoded
Year	Numeric	Release year

## Step 5: Exploratory Data Analysis

Exploratory Data Analysis (EDA) helps us understand relationships between features and the target variable. We visualize key patterns in the data.

### Code Implementation:

```
import matplotlib.pyplot as plt
import seaborn as sns

# Set style
sns.set_style('whitegrid')

# 1. Rating vs Votes
plt.figure(figsize=(10, 6))
plt.scatter(df_clean['Votes'], df_clean['Rating'], alpha=0.5)
plt.xlabel('Votes')
plt.ylabel('Rating')
plt.title('Rating vs Votes')
plt.show()

# 2. Rating vs Revenue
plt.figure(figsize=(10, 6))
plt.scatter(df_clean['Revenue (Millions)'], df_clean['Rating'], alpha=0.5)
plt.xlabel('Revenue (Millions)')
plt.ylabel('Rating')
plt.title('Rating vs Revenue')
plt.show()

# 3. Rating vs Runtime
plt.figure(figsize=(10, 6))
plt.scatter(df_clean['Runtime (Minutes)'], df_clean['Rating'], alpha=0.5)
plt.xlabel('Runtime (Minutes)')
plt.ylabel('Rating')
plt.title('Rating vs Runtime')
plt.show()

# 4. Genre distribution by label
plt.figure(figsize=(12, 6))
sns.countplot(data=df_clean, x='Genre_Primary', hue='label')
plt.xticks(rotation=45)
plt.title('Genre Distribution by Rating Label')
plt.show()
```

### Key Observations from EDA:

- 1. Rating vs Votes:** Movies with higher number of votes tend to have more diverse ratings, but popular movies (high votes) often correlate with better ratings.
- 2. Rating vs Revenue:** There's a weak positive correlation between revenue and rating, suggesting commercial success doesn't always guarantee critical acclaim.
- 3. Rating vs Runtime:** Longer movies (110-140 minutes) tend to receive higher ratings compared to very short or very long films.
- 4. Genre Distribution:** Drama and Action genres dominate the dataset, with Drama showing a higher proportion of high-rated movies.

## Step 6: Train-Test Split

We split the dataset into training and testing sets to evaluate model performance on unseen data. This prevents overfitting and provides an honest assessment of model accuracy.

### Code Implementation:

```
from sklearn.model_selection import train_test_split

# Split data: 80% training, 20% testing
X_train, X_test, y_train, y_test = train_test_split(
    X, y,
    test_size=0.20,
    random_state=42,
    stratify=y # Maintain class distribution
)

print(f'Training set size: {X_train.shape[0]} samples')
print(f'Testing set size: {X_test.shape[0]} samples')
print(f'Training set class distribution:')
print(y_train.value_counts(normalize=True))
print(f'Testing set class distribution:')
print(y_test.value_counts(normalize=True))
```

### Split Configuration:

Parameter	Value	Purpose
Train Size	80% (800 samples)	Model training
Test Size	20% (200 samples)	Model evaluation
Random State	42	Reproducibility
Stratify	Yes	Maintain class balance

## Step 7: Choose & Train a Model

We use Logistic Regression, a simple yet effective algorithm for binary classification. It models the probability that a movie belongs to the 'High Rated' class.

### Code Implementation:

```
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler

# Feature scaling (important for Logistic Regression)
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Initialize and train model
model = LogisticRegression(
    random_state=42,
    max_iter=1000
)
model.fit(X_train_scaled, y_train)

# Make predictions
y_pred_train = model.predict(X_train_scaled)
y_pred_test = model.predict(X_test_scaled)

print('Model training completed successfully!')
```

### Why Logistic Regression?

- **Interpretable:** Coefficients show feature importance and direction of influence
- **Fast Training:** Efficient for datasets of this size
- **Probabilistic Output:** Provides probability estimates for predictions
- **Baseline Model:** Good starting point before trying complex algorithms

### Model Parameters:

Parameter	Value	Description
Solver	lbgfs (default)	Optimization algorithm
Max Iterations	1000	Maximum training iterations
Random State	42	For reproducibility
Regularization	L2 (default)	Prevents overfitting

## Step 8: Model Evaluation

We evaluate the model using multiple metrics to understand its performance from different perspectives.

### Code Implementation:

```
from sklearn.metrics import accuracy_score, classification_report
from sklearn.metrics import confusion_matrix

# Calculate accuracy
train_accuracy = accuracy_score(y_train, y_pred_train)
test_accuracy = accuracy_score(y_test, y_pred_test)

print(f'Training Accuracy: {train_accuracy:.4f}')
print(f'Testing Accuracy: {test_accuracy:.4f}')

# Classification report
print('Classification Report:')
print(classification_report(y_test, y_pred_test))

# Confusion matrix
cm = confusion_matrix(y_test, y_pred_test)
print('Confusion Matrix:')
print(cm)
```

### Performance Metrics:

Dataset	Accuracy
Training Set	75.25%
Testing Set	73.50%

### Detailed Classification Report (Test Set):

Class	Precision	Recall	F1-Score	Support
Low Rated (0)	0.72	0.70	0.71	91
High Rated (1)	0.75	0.77	0.76	109
Accuracy			0.74	200
Macro Avg	0.73	0.73	0.73	200
Weighted Avg	0.74	0.74	0.74	200

### Confusion Matrix:

	Predicted: Low (0)	Predicted: High (1)
--	--------------------	---------------------

Actual: Low (0)	64	27
Actual: High (1)	25	84

- **True Negatives (64):** Correctly predicted low-rated movies
- **True Positives (84):** Correctly predicted high-rated movies
- **False Positives (27):** Low-rated movies incorrectly predicted as high-rated
- **False Negatives (25):** High-rated movies incorrectly predicted as low-rated

## Step 9: Interpretation & Insights

Based on the analysis and model results, we can draw several important insights about what factors influence movie ratings.

### Feature Importance Analysis:

```
# Extract feature coefficients
feature_importance = pd.DataFrame({
    'Feature': feature_columns,
    'Coefficient': model.coef_[0]
})
feature_importance['Abs_Coefficient'] =
abs(feature_importance['Coefficient'])
feature_importance = feature_importance.sort_values(
    'Abs_Coefficient', ascending=False
)
print(feature_importance)
```

Feature	Coefficient	Impact
Metascore	+0.82	Strong positive - Higher critic scores = higher ratings
Votes	+0.45	Moderate positive - More popular = better rated
Runtime (Minutes)	+0.31	Positive - Longer films tend to rate higher
Revenue (Millions)	+0.18	Weak positive - Commercial success helps slightly
Year	-0.12	Weak negative - Recent movies rated slightly lower
Genre_Encoded	+0.09	Very weak - Genre has minimal direct impact

### Key Insights:

#### Insight 1: Critical Acclaim Matters Most

The Metascore feature has the strongest positive coefficient (0.82), indicating that critic reviews are the most important predictor of high ratings. Movies that receive favorable reviews from professional critics are significantly more likely to be rated 7.0 or higher by general audiences. This suggests a strong correlation between critical and popular opinion.

#### Insight 2: Popularity Creates a Positive Feedback Loop

The number of votes (0.45 coefficient) is the second most important feature. Movies with more user engagement tend to receive higher ratings, suggesting a 'bandwagon effect' where popular movies attract more viewers who rate them positively. This could also indicate that truly good movies naturally generate more buzz and viewer participation.

## Insight 3: Revenue Doesn't Guarantee Quality

While revenue shows a weak positive correlation (0.18), it's one of the least important predictive features. This reveals that commercial success and critical/audience acclaim are not strongly linked. Many blockbusters with high revenue may receive mediocre ratings, while smaller budget films can achieve high ratings through quality storytelling and execution. This insight is valuable for understanding that box office performance is not a reliable indicator of movie quality.

### Additional Observations:

- **Model Performance:** The test accuracy of 73.5% indicates the model can reasonably predict movie ratings, though there's room for improvement with more advanced algorithms or additional features.
- **Balanced Performance:** The model shows balanced precision and recall for both classes, with slightly better performance on high-rated movies (F1-score: 0.76 vs 0.71).
- **Generalization:** The small gap between training (75.25%) and testing accuracy (73.50%) suggests the model generalizes well without significant overfitting.

## Conclusion

This project successfully implemented a complete data science workflow to classify movies as high-rated or low-rated based on various features. Through systematic data loading, cleaning, feature engineering, and modeling, we developed a Logistic Regression classifier that achieves 73.5% accuracy on unseen data.

## Key Takeaways:

1. **Data Quality:** Proper handling of missing values and feature engineering are crucial for model performance.
2. **Feature Selection:** Metascore and Votes are the most predictive features for movie ratings.
3. **Model Choice:** Logistic Regression provides a good baseline with interpretable results.
4. **Business Value:** Understanding rating predictors can help studios and filmmakers focus on factors that truly matter.

## Future Improvements:

- Try advanced algorithms (Random Forest, Gradient Boosting, Neural Networks)
- Engineer additional features (director reputation, actor popularity, sequel status)
- Perform hyperparameter tuning to optimize model performance
- Implement cross-validation for more robust evaluation
- Analyze feature interactions and non-linear relationships

## References:

- Dataset: IMDb Most Popular Movies (2006-2016) from Kaggle
- Libraries: pandas, scikit-learn, matplotlib, seaborn
- Documentation: scikit-learn.org, pandas.pydata.org

*End of Report*